

## Import Libraries

```
In [16]: import os
import torch
import torch.nn as nn
from torchsummary import summary
import torch.optim as optim
import torch.nn.functional as F
from torch.utils import data
from torchvision import datasets, transforms, models
from collections import Counter
from torchvision.transforms import Resize, CenterCrop, ToTensor, Normalize, GaussianBlur, Grayscale
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader, random_split, Subset, Dataset

import cv2
from PIL import Image, ImageFilter, ImageOps
from torchvision.transforms.functional import to_pil_image
from skimage.metrics import peak_signal_noise_ratio, structural_similarity
import matplotlib.pyplot as plt
import torch.optim as optim
from tqdm import tqdm

from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
from torch.optim import lr_scheduler

from tkinter import filedialog, Tk

from gradcam import GradCAM, GradCAMpp
from gradcam.utils import visualize_cam

import numpy as np
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_curve, auc

import dill
import torch.multiprocessing as mp

from captum.attr import IntegratedGradients
from captum.attr import visualization as viz
```

```
In [17]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu") # as M1 Mac Does not have a dedicated GPU
```

**Specify the path to the locally saved dataset**

```
In [18]: train_data_dir = '/Users/savin/Desktop/FYP/Implementation/kaggle_dataset/AugmentedAlzheimerDataset'
original_dataset = datasets.ImageFolder(train_data_dir)
```

```
In [19]: class ContrastStretching:
def __call__(self, img):
    # Convert PIL Image to NumPy array
    img_np = np.array(img)

    # Check if the image is grayscale or RGB
    if img_np.ndim == 2: # Grayscale image
        img_np = self.apply_contrast_stretching(img_np)
    elif img_np.ndim == 3: # RGB image
        # Apply contrast stretching to each channel individually
        for i in range(img_np.shape[-1]):
            img_np[:, :, i] = self.apply_contrast_stretching(img_np[:, :, i])

    # Convert back to PIL Image
    return Image.fromarray(img_np.astype('uint8'))

def apply_contrast_stretching(self, channel):
    in_min, in_max = np.percentile(channel, (0, 100))
    out_min, out_max = 0, 255
    channel = np.clip((channel - in_min) * (out_max - out_min) / (in_max - in_min) + out_min, out_min, out_max)
    return channel

class UnsharpMask:
def __init__(self, radius=1, percent=100, threshold=3):
    self.radius = radius
    self.percent = percent
    self.threshold = threshold

def __call__(self, img):
    return img.filter(ImageFilter.UnsharpMask(radius=self.radius,
                                              percent=self.percent,
                                              threshold=self.threshold
                                              ))

class GaussianBlur:
def __init__(self, kernel_size, sigma=(0.1, 2.0)):
    self.kernel_size = kernel_size
    self.sigma = sigma

def __call__(self, img):
    sigma = np.random.uniform(self.sigma[0], self.sigma[1])
    img = img.filter(ImageFilter.GaussianBlur(sigma))
    return img
```

#### Preprocess the dataset

```
In [20]: preprocess_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    ContrastStretching(),
    UnsharpMask(radius=1, percent=100, threshold=3),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 0.5)),
    transforms.ToTensor()
])

preprocessed_dataset = datasets.ImageFolder(root=train_data_dir, transform=preprocess_transform)

data_loader = DataLoader(preprocessed_dataset, batch_size=32, shuffle=True)
```

```
In [21]: def calculate_psnr_ssim(original_dataset, preprocessed_dataset, num_samples=100, resize=(224, 224)):
    psnr_values = []
    ssim_values = []

    for i in range(num_samples):
        original_img = original_dataset[i][0] # original MRI
        preprocessed_img = preprocessed_dataset[i][0] # preprocessed MRI

        if not isinstance(original_img, Image.Image):
            original_img = to_pil_image(original_img)
        if not isinstance(preprocessed_img, Image.Image):
            preprocessed_img = to_pil_image(preprocessed_img)

        # Convert MRI to grayscale
        original_img = original_img.convert("L")
        preprocessed_img = preprocessed_img.convert("L")

        # Resize images
        original_img = original_img.resize(resize)
        preprocessed_img = preprocessed_img.resize(resize)

        # Convert images to numpy arrays
        original_img_np = np.array(original_img)
        preprocessed_img_np = np.array(preprocessed_img)

        # Calculate PSNR and SSIM
        psnr = peak_signal_noise_ratio(original_img_np, preprocessed_img_np, data_range=original_img_np.max() -
        ssim = structural_similarity(original_img_np, preprocessed_img_np)

        psnr_values.append(psnr)
        ssim_values.append(ssim)

    # Compute average PSNR and SSIM
    avg_psnr = np.mean(psnr_values)
    avg_ssim = np.mean(ssim_values)

    return avg_psnr, avg_ssim

# Example usage
avg_psnr, avg_ssim = calculate_psnr_ssim(original_dataset, preprocessed_dataset)
print(f"Average PSNR: {avg_psnr}, Average SSIM: {avg_ssim}")
```

Average PSNR: 27.207944599068227, Average SSIM: 0.9001335626106904

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 2.0)),
    transforms.Lambda(lambda x: x.filter(ImageFilter.UnsharpMask(radius=2, percent=150, threshold=3))),
    transforms.ToTensor()
)
```

Average PSNR: 9.729432125669954, Average SSIM: 0.2833625979364462

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    transforms.ToTensor()
)
```

Average PSNR: 10.2026194786185, Average SSIM: 0.32374658927511385

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 2.0)),
    transforms.ToTensor(),
)
```

Average PSNR: 34.479706650199184, Average SSIM: 0.9638484917028203

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 1.0)),
    transforms.ToTensor(),
)
```

Average PSNR: 40.03000513450271, Average SSIM: 0.9923509776637894

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    transforms.Lambda(lambda img: img.filter(ImageFilter.UnsharpMask(radius=2, percent=100, threshold=3))),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 0.5)),
    transforms.ToTensor()
)
```

Average PSNR: 28.919164968402907, Average SSIM: 0.9646276430637585

```
preprocess_transform = transforms.Compose(
    transforms.Resize((224, 224)),
    ContrastStretching(),
    transforms.Lambda(lambda img: img.filter(ImageFilter.UnsharpMask(radius=1, percent=100, threshold=3))),
    GaussianBlur(kernel_size=(5, 5), sigma=(0.1, 0.5)),
    transforms.ToTensor()
)
```

Average PSNR: 27.194490517269728, Average SSIM: 0.8948121010151182

#### MRI scan counts in each class of the dataset

```
In [23]: ted = '/Users/savin/Desktop/FYP/Implementation/kaggle_dataset/AugmentedAlzheimerDataset/MildDemented'
emented = '/Users/savin/Desktop/FYP/Implementation/kaggle_dataset/AugmentedAlzheimerDataset/ModerateDemented'
ed = '/Users/savin/Desktop/FYP/Implementation/kaggle_dataset/AugmentedAlzheimerDataset/NonDemented'
emented = '/Users/savin/Desktop/FYP/Implementation/kaggle_dataset/AugmentedAlzheimerDataset/VeryMildDemented'

dDemented = len(os.listdir(MildDemented))
erateDemented = len(os.listdir(ModerateDemented))
Demented = len(os.listdir(NonDemented))
yMildDemented = len(os.listdir(VeryMildDemented))

umber of images in MildDemented: {count_MildDemented}")
umber of images in ModerateDemented: {count_ModerateDemented}")
umber of images in NonDemented: {count_NonDemented}")
umber of images in VeryMildDemented: {count_VeryMildDemented}")

nTotal MRIs in the dataset = {count_MildDemented+count_ModerateDemented+count_NonDemented+count_VeryMildDemented}
```

```
Number of images in MildDemented: 8960
Number of images in ModerateDemented: 6464
Number of images in NonDemented: 9600
Number of images in VeryMildDemented: 8960
```

Total MRIs in the dataset = 33984

#### Sample MRI before and after preprocessing

```
In [22]: sample_image_path = os.path.join(MildDemented, os.listdir(MildDemented)[0])

original_image = Image.open(sample_image_path)

# Apply the preprocessing transforms
preprocessed_image = preprocess_transform(original_image)


preprocessed_image = transforms.ToPILImage()(preprocessed_image)

# Display the images
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
plt.imshow(original_image)
plt.title("Original Image")
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(preprocessed_image)
plt.title("Preprocessed Image")
plt.axis('off')

plt.show()
```



### Dataset splitting & creating DataLoaders

```
In [24]: train_size = int(0.70 * len(preprocessed_dataset))
val_size = int(0.15 * len(preprocessed_dataset))
test_size = len(preprocessed_dataset) - train_size - val_size

train_dataset, val_dataset, test_dataset = random_split(preprocessed_dataset, [train_size, val_size, test_size])

# Create DataLoaders
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=16, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False)

dataloaders = {'train': train_loader, 'val': val_loader, 'test': test_loader}
```

```
In [40]: all_labels = [label for _, label in train_dataset]
class_distribution = Counter(all_labels)
print(class_distribution)
```

```
Counter({2: 6737, 0: 6238, 3: 6236, 1: 4577})
```

## Building CNN Model 2

```
In [25]: class SEBlock(nn.Module):
    def __init__(self, in_channels, reduction=32):
        super(SEBlock, self).__init__()
        self.avg_pool = nn.AdaptiveAvgPool2d(1)
        self.fc = nn.Sequential(
            nn.Linear(in_channels, in_channels // reduction, bias=False),
            nn.ReLU(inplace=True),
            nn.Linear(in_channels // reduction, in_channels, bias=False),
            nn.Sigmoid()
        )

    def forward(self, x):
        b, c, _, _ = x.size()
        y = self.avg_pool(x).view(b, c)
        y = self.fc(y).view(b, c, 1, 1)
        return x * y.expand_as(x)

class ResidualBlock(nn.Module):
    def __init__(self, in_channels):
        super(ResidualBlock, self).__init__()
        self.conv = nn.Conv2d(in_channels, in_channels, kernel_size=3, padding=1, bias=False)
        self.bn = nn.BatchNorm2d(in_channels)

    def forward(self, x):
        residual = x
        out = F.relu(self.bn(self.conv(x)))
        out += residual
        return F.relu(out)

class CustomEfficientNet(nn.Module):
    def __init__(self, num_classes=4):
        super(CustomEfficientNet, self).__init__()
        self.base_model = models.efficientnet_b0(pretrained=True)

        for param in self.base_model.parameters():
            param.requires_grad = False

        # Replace the classifier with a new one
        num_fts = self.base_model.classifier[1].in_features
        self.classifier = nn.Sequential(
            nn.Dropout(0.2),
            nn.Linear(num_fts, 512),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(512, num_classes),
        )

        # Add SEBlock and ResidualBlock to the end of the features, before the avgpool
        self.base_model.features.add_module("SEBlock", SEBlock(1280)) # Adjust in_channels accordingly

    def forward(self, x):
        # Process through EfficientNet up to before avgpool
        x = self.base_model.features(x)

        # Now x is a 4D tensor, and we can apply SEBlock and ResidualBlock directly
        # No need for separate calls, as they are part of the features module

        # Apply avgpool and classifier
        x = self.base_model.avgpool(x)
        x = torch.flatten(x, 1)
        x = self.classifier(x)
        return x

# Creating the model and moving to device
model2 = CustomEfficientNet(num_classes=4).to(device)

# For a summary, ensure the input size matches your dataset
summary(model2, (3, 224, 224))
```



**Train Customized EfficientNet-B0**



```

In [26]: criterion = nn.CrossEntropyLoss()
trainable_params = filter(lambda p: p.requires_grad, model2.parameters())
optimizer = torch.optim.Adam(trainable_params, lr=0.001)
scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
torch.autograd.set_detect_anomaly(True)

def train_epoch(epoch_index, train_loader, model, optimizer):
    model.train()
    running_loss = 0.0
    correct_pred = 0
    total_pred = 0

    for inputs, labels in tqdm(train_loader, desc=f"Epoch {epoch_index + 1} [train] Progress", unit="batch"):
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predictions = torch.max(outputs, 1)
        correct_pred += (predictions == labels).sum().item()
        total_pred += labels.size(0)

    avg_loss = running_loss / len(train_loader)
    avg_acc = correct_pred / total_pred
    print(f'train Loss: {avg_loss:.4f} Acc: {avg_acc:.4f}')
    return avg_loss, avg_acc

def validate_epoch(epoch_index, val_loader, model):
    model.eval()
    running_loss = 0.0
    correct_pred = 0
    total_pred = 0

    with torch.no_grad():
        for inputs, labels in tqdm(val_loader, desc=f"Epoch {epoch_index + 1} [val] Progress", unit="batch"):
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)

            running_loss += loss.item()
            _, predictions = torch.max(outputs, 1)
            correct_pred += (predictions == labels).sum().item()
            total_pred += labels.size(0)

    avg_loss = running_loss / len(val_loader)
    avg_acc = correct_pred / total_pred
    print(f'val Loss: {avg_loss:.4f} Acc: {avg_acc:.4f}')
    return avg_loss, avg_acc

# Training loop
num_epochs = 100
train_losses, train_accuracies = [], []
val_losses, val_accuracies = [], []
best_val_loss = float('inf')
patience = 8

for epoch in range(num_epochs):
    train_loss, train_acc = train_epoch(epoch, train_loader, model2, optimizer)
    val_loss, val_acc = validate_epoch(epoch, val_loader, model2)

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        trigger_times = 0
        torch.save(model2.state_dict(), 'model2_test1.pth')
    else:
        trigger_times += 1
        if trigger_times >= patience:
            print(f"Early stopping at epoch {epoch+1}")
            break


    train_losses.append(train_loss)
    train_accuracies.append(train_acc)
    val_losses.append(val_loss)
    val_accuracies.append(val_acc)

# Plotting
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.legend()
plt.title('Loss over epochs')

plt.subplot(1, 2, 2)
plt.plot(train_accuracies, label='Training Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.legend()
plt.title('Accuracy over epochs')

```

plt.show()



Classification Report of the trained Modified EfficientNet B0

```
In [27]: # classification report (val loader)
model2 = CustomEfficientNet(num_classes=4).to(device)

model2.load_state_dict(torch.load('model2_test1.pth'))

def evaluate_model(model, dataloader):
    model.eval()
    true_labels = []
    predictions = []

    with torch.no_grad():
        for inputs, labels in dataloader:
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)

            true_labels.extend(labels.cpu().numpy())
            predictions.extend(preds.cpu().numpy())

    return true_labels, predictions

# Evaluate the model
true_labels, predictions = evaluate_model(model2, val_loader)

# Print classification report
print(classification_report(true_labels, predictions, target_names=['MildDemented', 'ModerateDemented', 'NonDeme
```

	precision	recall	f1-score	support
MildDemented	0.71	0.73	0.72	1334
ModerateDemented	0.89	0.95	0.92	944
NonDemented	0.78	0.62	0.69	1469
VeryMildDemented	0.58	0.67	0.63	1350
accuracy			0.72	5097
macro avg	0.74	0.74	0.74	5097
weighted avg	0.73	0.72	0.72	5097

```
In [30]: # Save the best model locally (model2 - modified EfficientNetB0)

model_save_path = '/Users/savin/Desktop/FYP/Implementation'
os.makedirs(model_save_path, exist_ok=True)
model_save_file = os.path.join(model_save_path, 'model2_test2.pth')

torch.save(model2.state_dict(), model_save_file)

print(f'Model saved to {model_save_file}')
```

Model saved to /Users/savin/Desktop/FYP/Implementation/model2\_test2.pth

```
In [28]: def evaluate_model(model, dataloader):
        model.eval()
        true_labels = []
        predictions = []

        with torch.no_grad():
            for inputs, labels in dataloader:
                inputs, labels = inputs.to(device), labels.to(device)

                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)

                true_labels.extend(labels.cpu().numpy())
                predictions.extend(preds.cpu().numpy())

        return true_labels, predictions

# Evaluate the model
true_labels, predictions = evaluate_model(model2, test_loader)

# Print classification report
print(classification_report(true_labels, predictions, target_names=['MildDemented', 'ModerateDemented', 'NonDeme
```

	precision	recall	f1-score	support
MildDemented	0.68	0.73	0.71	1280
ModerateDemented	0.91	0.95	0.93	936
NonDemented	0.78	0.59	0.67	1508
VeryMildDemented	0.56	0.66	0.60	1375
accuracy			0.71	5099
macro avg	0.73	0.73	0.73	5099
weighted avg	0.72	0.71	0.71	5099

```
In [29]: # Confusion Matrix on Test Loader - modified EfficientNet B0

def get_predictions(model, dataloader):
    model.eval()
    true_labels = []
    predictions = []

    with torch.no_grad():
        for inputs, labels in dataloader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)


            true_labels.extend(labels.cpu().numpy())
            predictions.extend(preds.cpu().numpy())

    return true_labels, predictions

# Evaluate the model
true_labels, predictions = get_predictions(model2, test_loader)

# Compute the confusion matrix
cm = confusion_matrix(true_labels, predictions)
class_names = ['MildDemented', 'ModerateDemented', 'NonDemented', 'VeryMildDemented']

# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix (Test data)')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()
```



In [ ]: